

# PPOL 628: Text as Data – Computational Linguistics for Social Scientists

Class 12: Word Embeddings

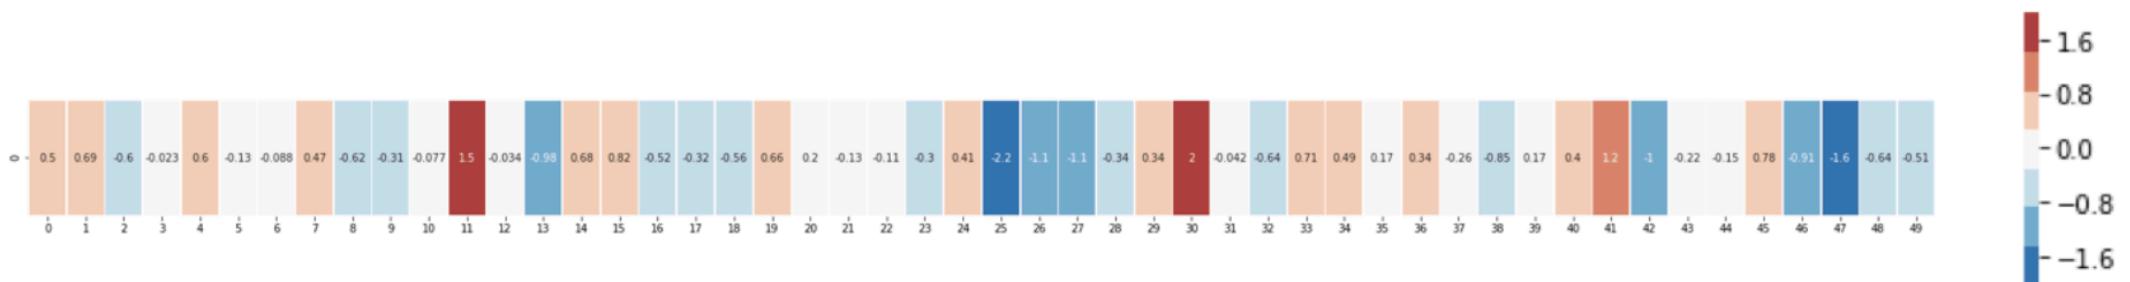
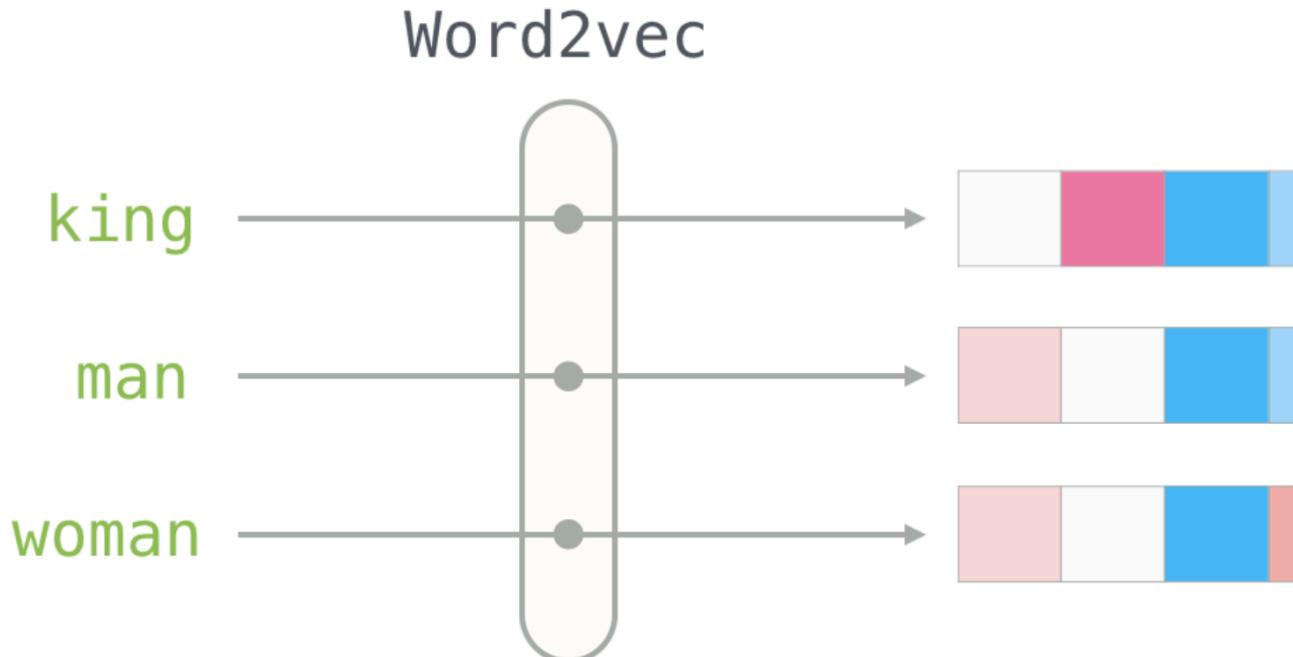
# Today

- Lecture: Word Embeddings.
  - Introduction to word embeddings.
  - Fitting Word2Vec.
  - Potential social science applications.
- Lab: topic\_models\_2.R
- Website: [github.com/matthewjdenny/PPOL\\_628\\_Text\\_As\\_Data](https://github.com/matthewjdenny/PPOL_628_Text_As_Data)

# Word Embeddings

- Humans have learned understand words as connected to meanings, and are able to think of synonyms, make analogies, etc.
  - In order for computers to mimic this type of behavior, we have to represent words in a way that computers can “understand”.
- One approach is to represent them as points in a (50-1000 dimensional) vector space. Meaning of words mapped to dimensions of vector space.
  - Turns out this approach works well in lots of NLP tasks – social science is still searching for a homerun application.
- Many papers + tutorials (which I screenshot from liberally in these slides):
  - Visualization: <http://jalammar.github.io/illustrated-word2vec/>
  - Jurafsky Slides: <https://web.stanford.edu/~jurafsky/slp3/slides/vector2.pdf>
  - Tutorial: [http://u.cs.biu.ac.il/~gonenhi/slides/WiDS\\_tutorial.pdf](http://u.cs.biu.ac.il/~gonenhi/slides/WiDS_tutorial.pdf)
  - GloVe Website: <https://nlp.stanford.edu/projects/glove/>

# Words Represented as Numerical Vectors

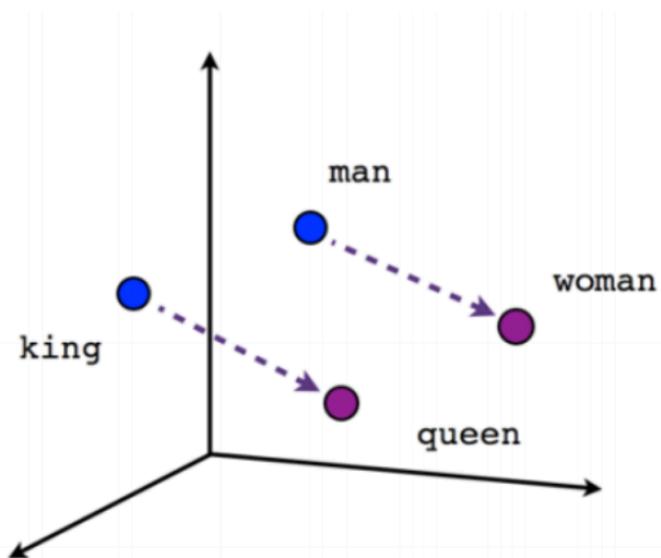


# Word Embeddings

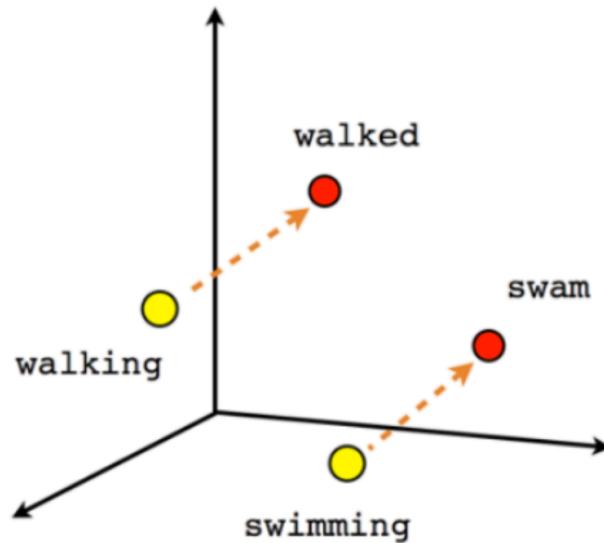
- There has been a huge amount of research into word embeddings or “vector space representations”, particularly as input for deep learning NLP models.
  - Literally thousands of academic papers.
  - Typically used as a form of preprocessing for deep learning algorithms.
  - Allows individual words to have some “intrinsic” meaning.
- Example: First ten dimensions of embedding for the word “football” in the GloVe 200 dimensional embedding space:
  - {-0.11, -0.34, -0.99, 0.85, 0.1, 0.74, -0.1, 0.36, 1.2, ...}

# Representing Relationships Between Words

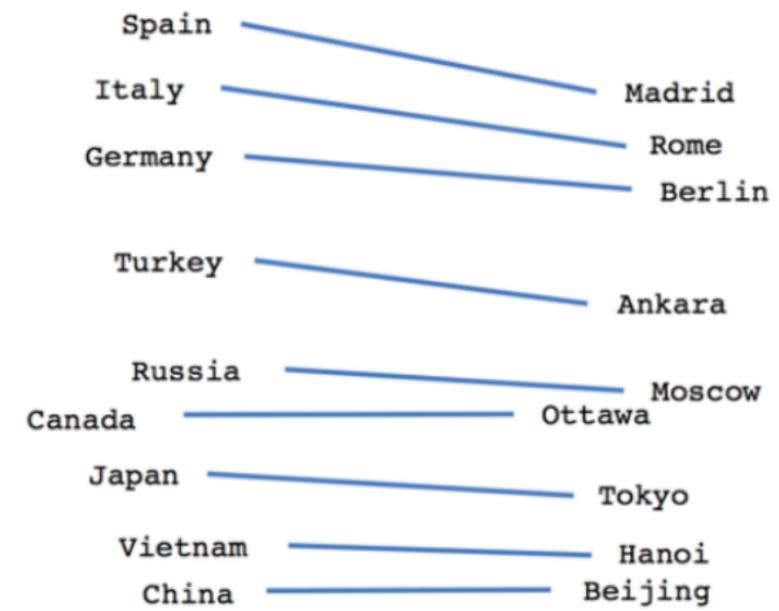
- Relative position of words in vector or embedding space can tell us about relationships among words.



Male-Female

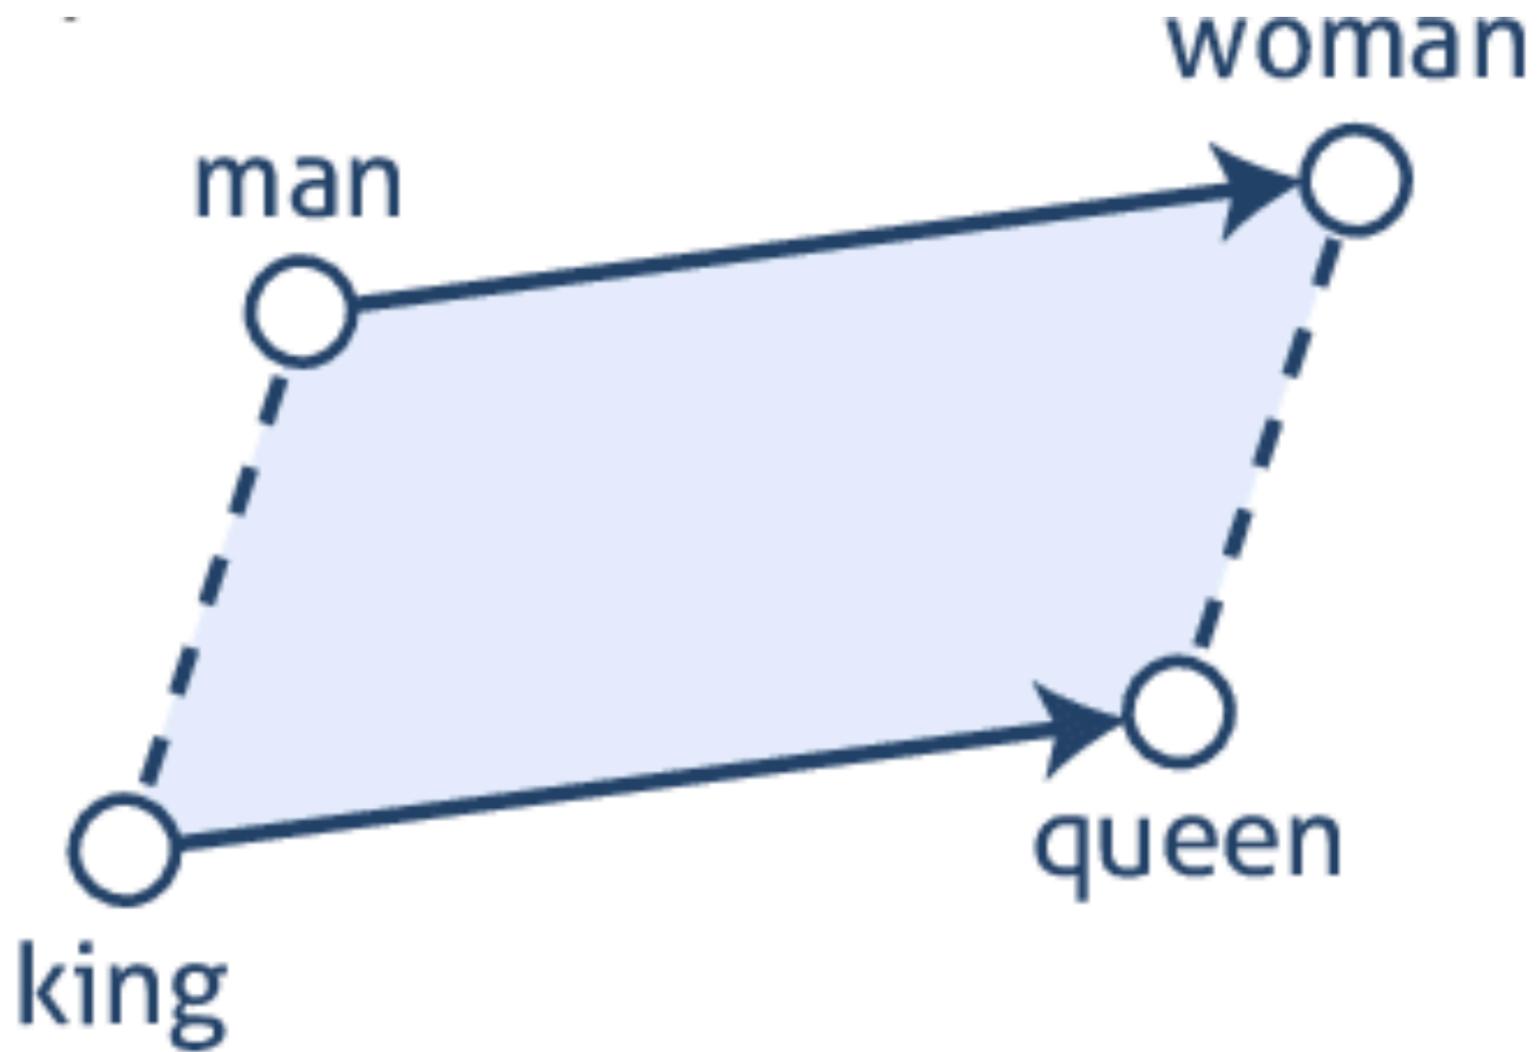


Verb tense



Country-Capital

# Solving Analogies



# Representation Learning

- Word embeddings are cool, and useful in a number of settings, but how do we generate them?
  - Answers come from field of “representation learning” where we use a machine learning model to learn the vector space representation of each word.
- Common models include Google’s Word2Vec, Stanford’s GloVe, Fasttext, etc. We will focus on the Word2Vec model in this lecture.
  - Goal of all of these models is to take a very large collection of documents as input and return a  $k$ -dimensional output numerical representation of each of the words in that corpus.
  - To have good performance, these models rely on absolutely enormous training corpora.
  - Typically have a neural network model architecture.

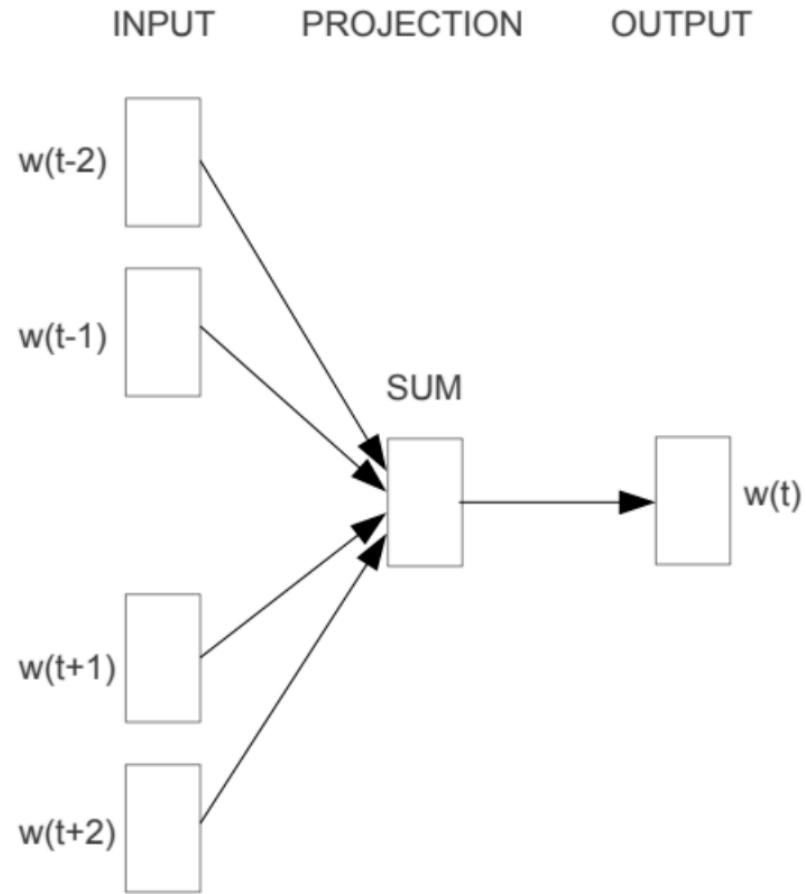
# Word2Vec

- Basic intuition: train a binary classifier to predict words that will show up close to each other. Use the weights the classifier learns for this task as the vector space embedding for each word.
- Magic advantage here is we do not need any human labeled training data -- we can just use the co-occurrence of terms near each other in real documents as our “training data”.
- Relatively fast to train – but is dependent on having lots of data.
- We will cover: "skip-gram with negative sampling" (SGNS)

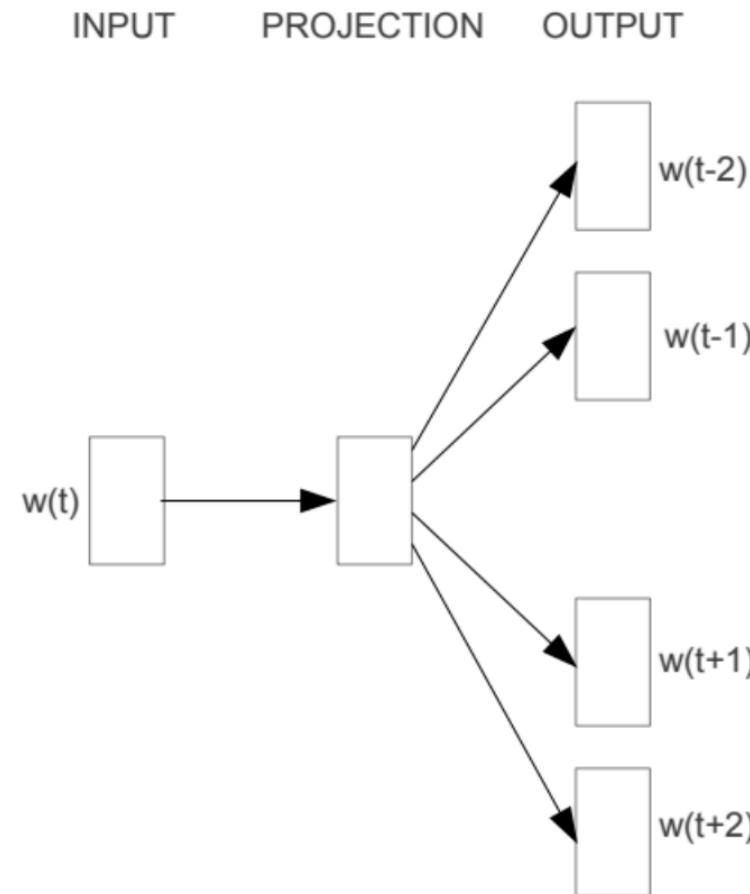
# Classification Tasks

- Prediction task is to predict what word will show up at a given ***target*** position ( $t$ ) in a document given some ***context*** words.
  - Two different ways to conceptualize the problem.
- Continuous Bag of Words (CBOW)
  - Use words that occur near the target word to predict the target word.
- Skip-Grams (SG)
  - Use the target word to predict the context words around that target.
  - Tends to provide better performance in practice.

# Classification Tasks



**CBOW**



**Skip-gram**

# Context Window

■ : Center Word

■ : Context Word

c=0    The cute cat jumps over the lazy dog.

c=1    The cute cat jumps over the lazy dog.

c=2    The cute cat jumps over the lazy dog.

# Skip-gram training data

- Typically break up the input training data into sentences.
  - Words that are close to other words but separated by sentence boundary can have little to no similarity.
  - Paragraphs with different topics, etc.

Training sentence:

... lemon, a **tablespoon of apricot jam** a pinch ...

c1

c2 target c3 c4

# Goal with Skip-gram Approach

Given a tuple (t,c) = target, context

- (*apricot*, *jam*)
- (*apricot*, *aardvark*)

Return probability that c is a real context word:

$$P(+ | t, c)$$

$$P(- | t, c) = 1 - P(+ | t, c)$$

# The Dot Product

- Consider two vectors:  $\mathbf{a} = [a_1, a_2, \dots, a_n]$  and  $\mathbf{b} = [b_1, b_2, \dots, b_n]$
- Then the dot product of those two vectors is:

$$\mathbf{a} \cdot \mathbf{b} = \sum_{i=1}^n a_i b_i = a_1 b_1 + a_2 b_2 + \cdots + a_n b_n$$

- Can serve as a measure of similarity between two vectors as it is inherently related to cosine similarity.

$$\cos(\theta) = \frac{\mathbf{u} \cdot \mathbf{v}}{\|\mathbf{u}\| \|\mathbf{v}\|}$$

# How to Generate These Probabilities?

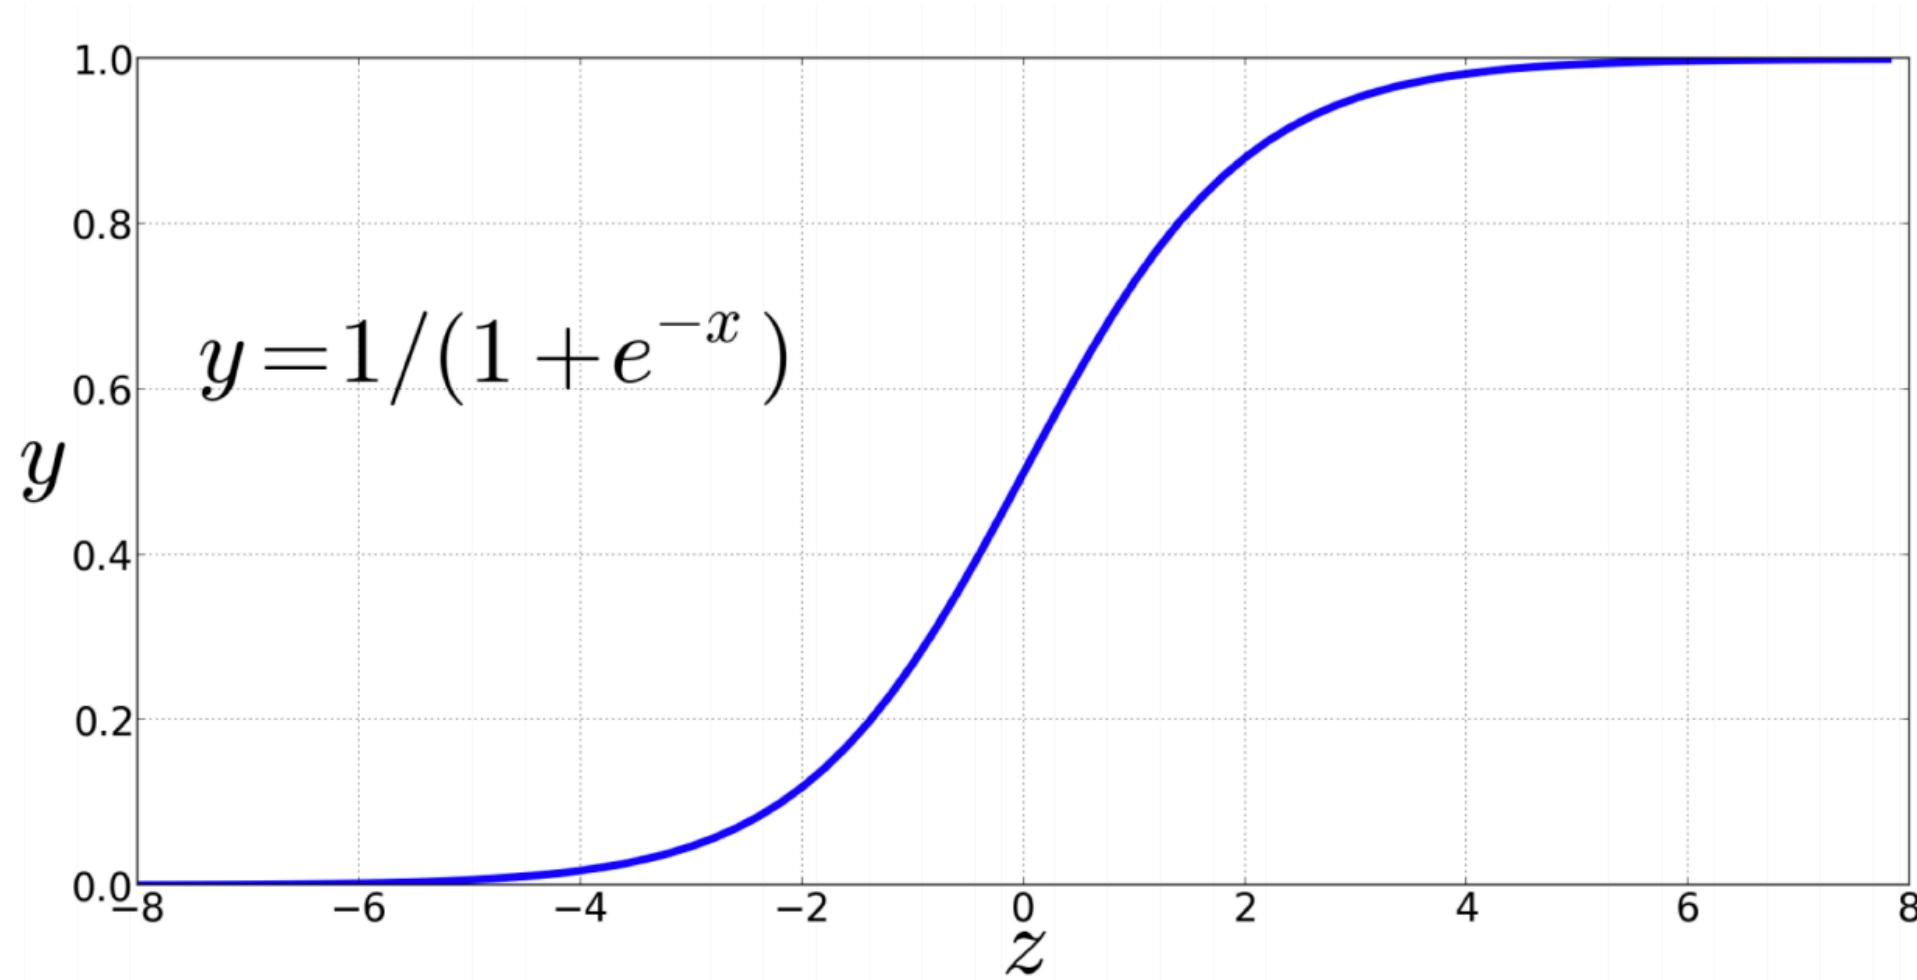
Intuition:

- Words are likely to appear near similar words
- Model similarity with dot-product!
- $\text{Similarity}(t,c) \propto t \cdot c$

Problem:

- *Dot product is not a probability!*
- *(Neither is cosine)*

# Logistic Sigmoid Function Gives us Probabilities



# Probabilities

- Can generate probabilities of a context word being a real context word for a given target by transforming dot product of context word co-occurring with target word in given context window across all times target word appears.
- Can roll up this probability over all words in context window.

$$P(+|t, c) = \frac{1}{1 + e^{-t \cdot c}}$$

$$\begin{aligned} P(-|t, c) &= 1 - P(+|t, c) \\ &= \frac{e^{-t \cdot c}}{1 + e^{-t \cdot c}} \end{aligned}$$

$$P(+|t, c_{1:k}) = \prod_{i=1}^k \frac{1}{1 + e^{-t \cdot c_i}}$$

$$\log P(+|t, c_{1:k}) = \sum_{i=1}^k \log \frac{1}{1 + e^{-t \cdot c_i}}$$

# Generating Positive and Negative Examples

Training sentence:

... lemon, a **tablespoon of apricot jam** a pinch ...

c1            c2    t            c3    c4

**positive examples +**

t            c

---

apricot tablespoon

apricot of

apricot preserves

apricot or

- For each positive example, we'll create  $k$  negative examples.
- Using *noise* words
- Any random word that isn't  $t$

# Generating Positive and Negative Examples

Training sentence:

... lemon, a **tablespoon** of **apricot** jam a pinch ...

c1

c2 t

c3 c4

**positive examples +**

t c

---

apricot tablespoon

apricot of

apricot preserves

apricot or

**negative examples - <sup>k=2</sup>**

t c t c

---

apricot aardvark apricot twelve

apricot puddle apricot hello

apricot where apricot dear

apricot coaxial apricot forever

# Picking Negative Examples

Could pick  $w$  according to their unigram frequency  $P(w)$

More common to chosen then according to  $p_\alpha(w)$

$$P_\alpha(w) = \frac{\text{count}(w)^\alpha}{\sum_w \text{count}(w)^\alpha}$$

$\alpha = \frac{3}{4}$  works well because it gives rare noise words slightly higher probability

To show this, imagine two events  $p(a) = .99$  and  $p(b) = .01$ :

$$P_\alpha(a) = \frac{.99^{.75}}{.99^{.75} + .01^{.75}} = .97$$

$$P_\alpha(b) = \frac{.01^{.75}}{.99^{.75} + .01^{.75}} = .03$$

# Model Training Objective

Let's represent words as vectors of some length (say 300), randomly initialized.

So we start with  $300 * V$  random parameters

Over the entire training set, we'd like to adjust those word vectors such that we

- Maximize the similarity of the **target word, context word** pairs  $(t,c)$  drawn from the positive data
- Minimize the similarity of the  $(t,c)$  pairs drawn from the negative data.

# Model Training Process

Iterative process.

We'll start with 0 or random weights

Then adjust the word weights to

- make the positive pairs more likely
- and the negative pairs less likely

over the entire training set:

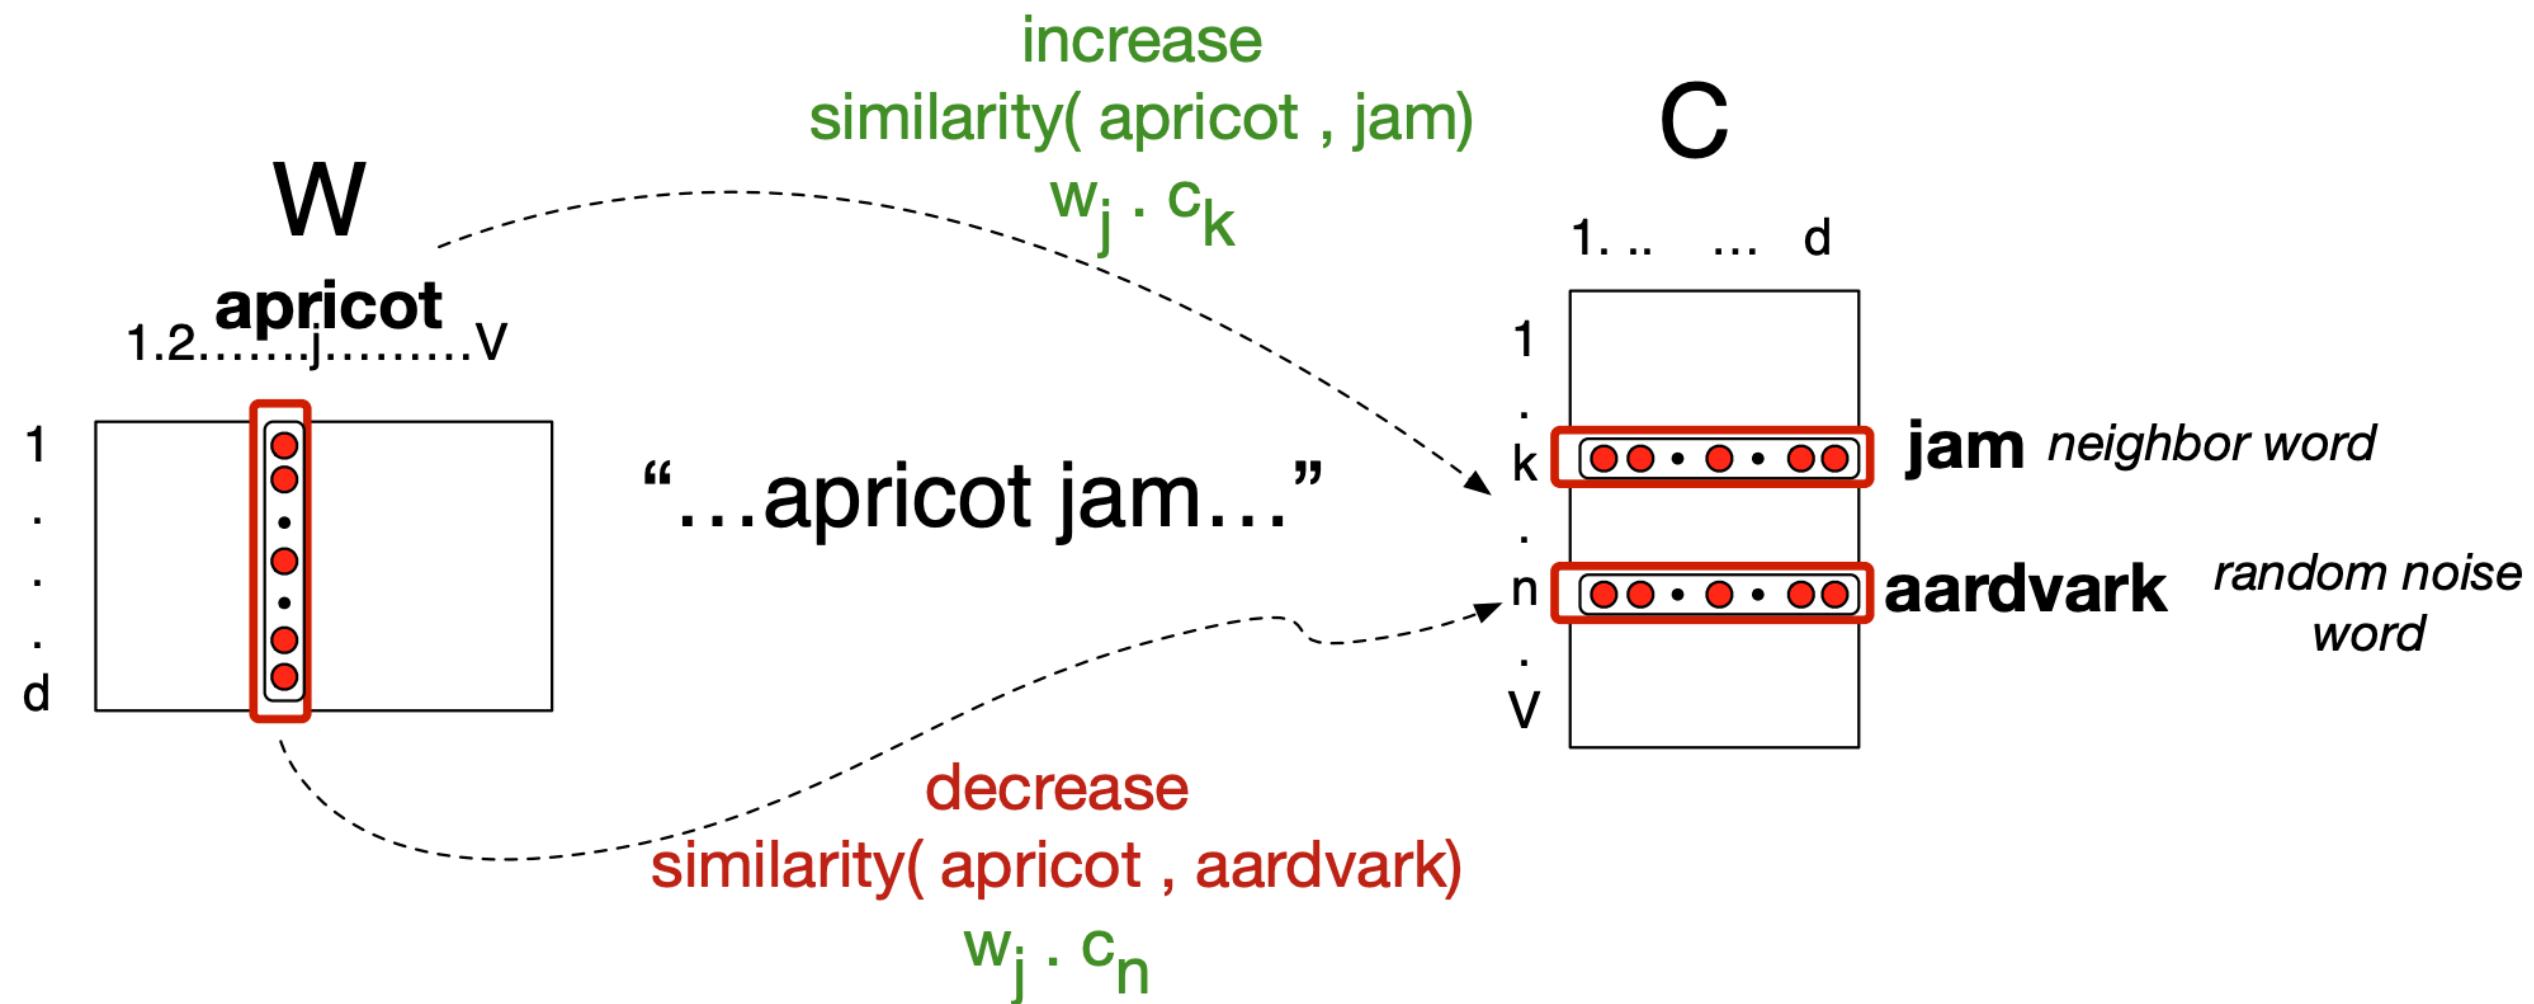
# Objective Function

We want to maximize...

$$\sum_{(t,c) \in +} \log P(+|t, c) + \sum_{(t,c) \in -} \log P(-|t, c)$$

Maximize the + label for the pairs from the positive training data, and the – label for the pairs sample from the negative data.

# Example of Iterative Adjustment



# Overview of Word2Vec Training

Start with  $V$  random 300-dimensional vectors as initial embeddings

Use logistic regression, the second most basic classifier used in machine learning after naïve bayes

- Take a corpus and take pairs of words that co-occur as positive examples
- Take pairs of words that don't co-occur as negative examples
- Train the classifier to distinguish these by slowly adjusting all the embeddings to improve the classifier performance
- Throw away the classifier code and keep the embeddings.

# Context Window Effects Results

- Changing size of context window for training will lead to model identifying different terms as being most similar.

$C = \pm 2$  The nearest words to *Hogwarts*:

- *Sunnydale*
- *Evernight*

$C = \pm 5$  The nearest words to *Hogwarts*:

- *Dumbledore*
- *Malfoy*
- *halfblood*

# Using Word Embeddings

- Now that we have numbers represented as embedding vectors, we can apply lots of nice mathematical concepts:
  - Similarity between words as distance in the embedding space.
  - Relationships between words as a ray (a direction in the embedding space + a distance)
  - Geometry of concept relations: e.g. analogies.
  - Calibrated embeddings as a way to translate text, concepts.
  - Find different words that share the same meaning when used in different contexts (e.g. legislators and CEOs use different terms to talk about the same regulations).
  - Look at changes in term associations over time, by covariates, to uncover changing word use, linguistic bias, etc.
    - Requires you to train your own embedding model.

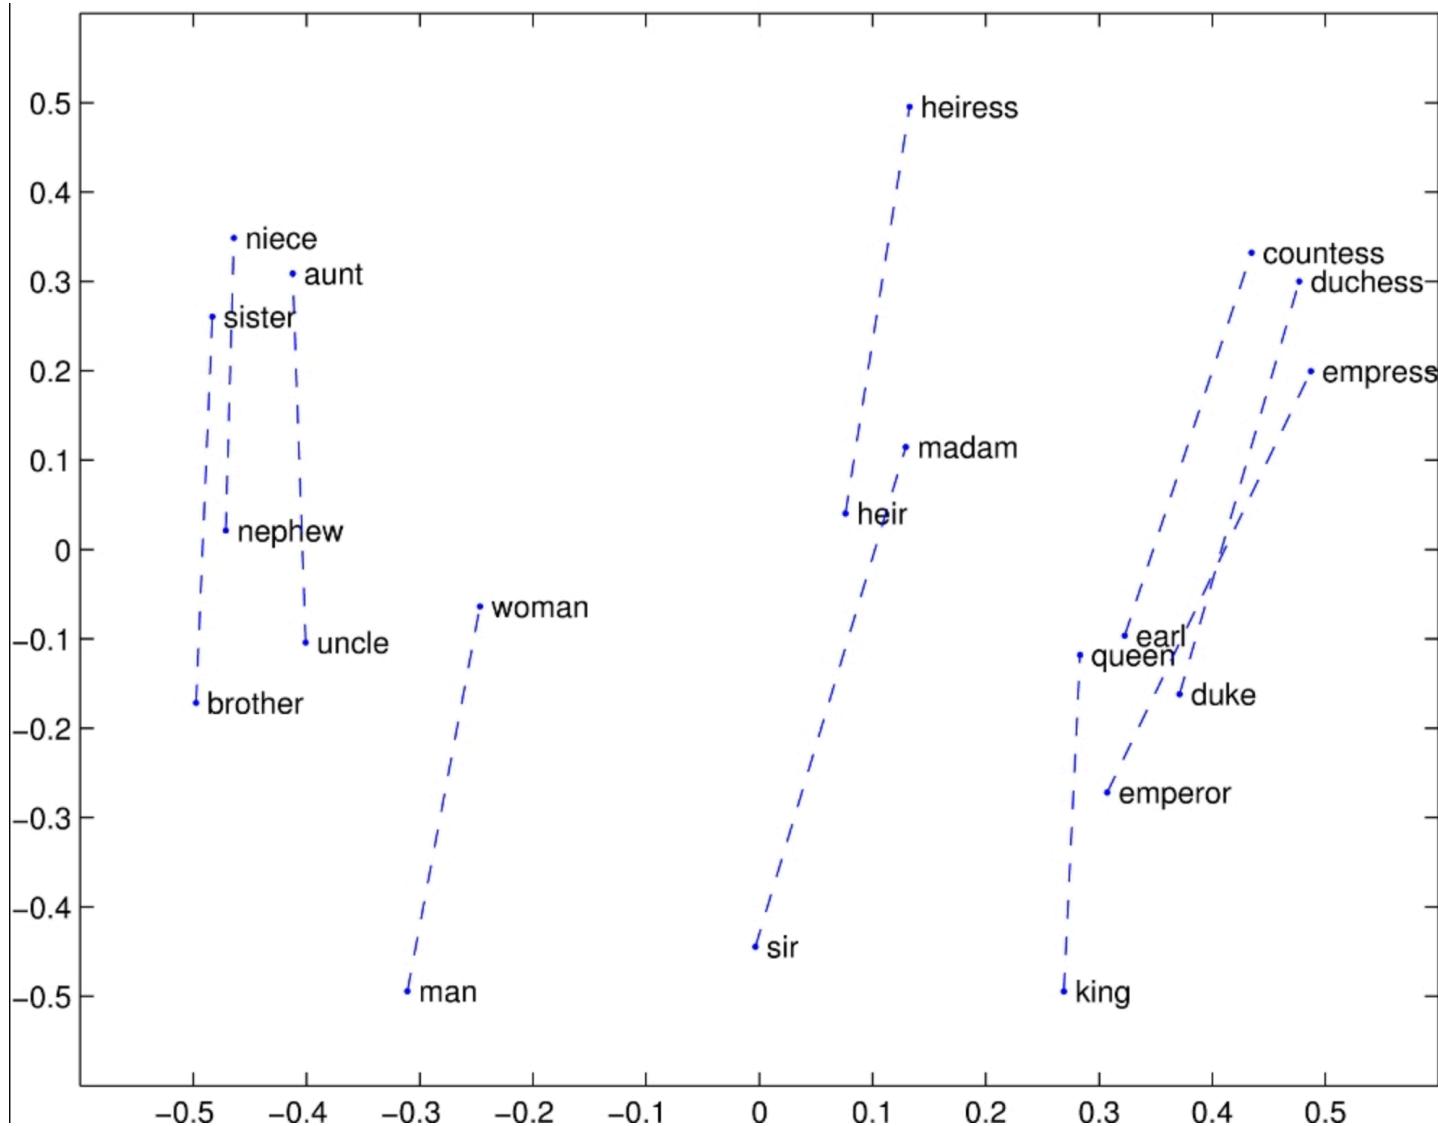
# Cosine Similarity

- Input: Let  $\mathbf{A}$  and  $\mathbf{B}$  be term vectors as with Euclidean Distance.

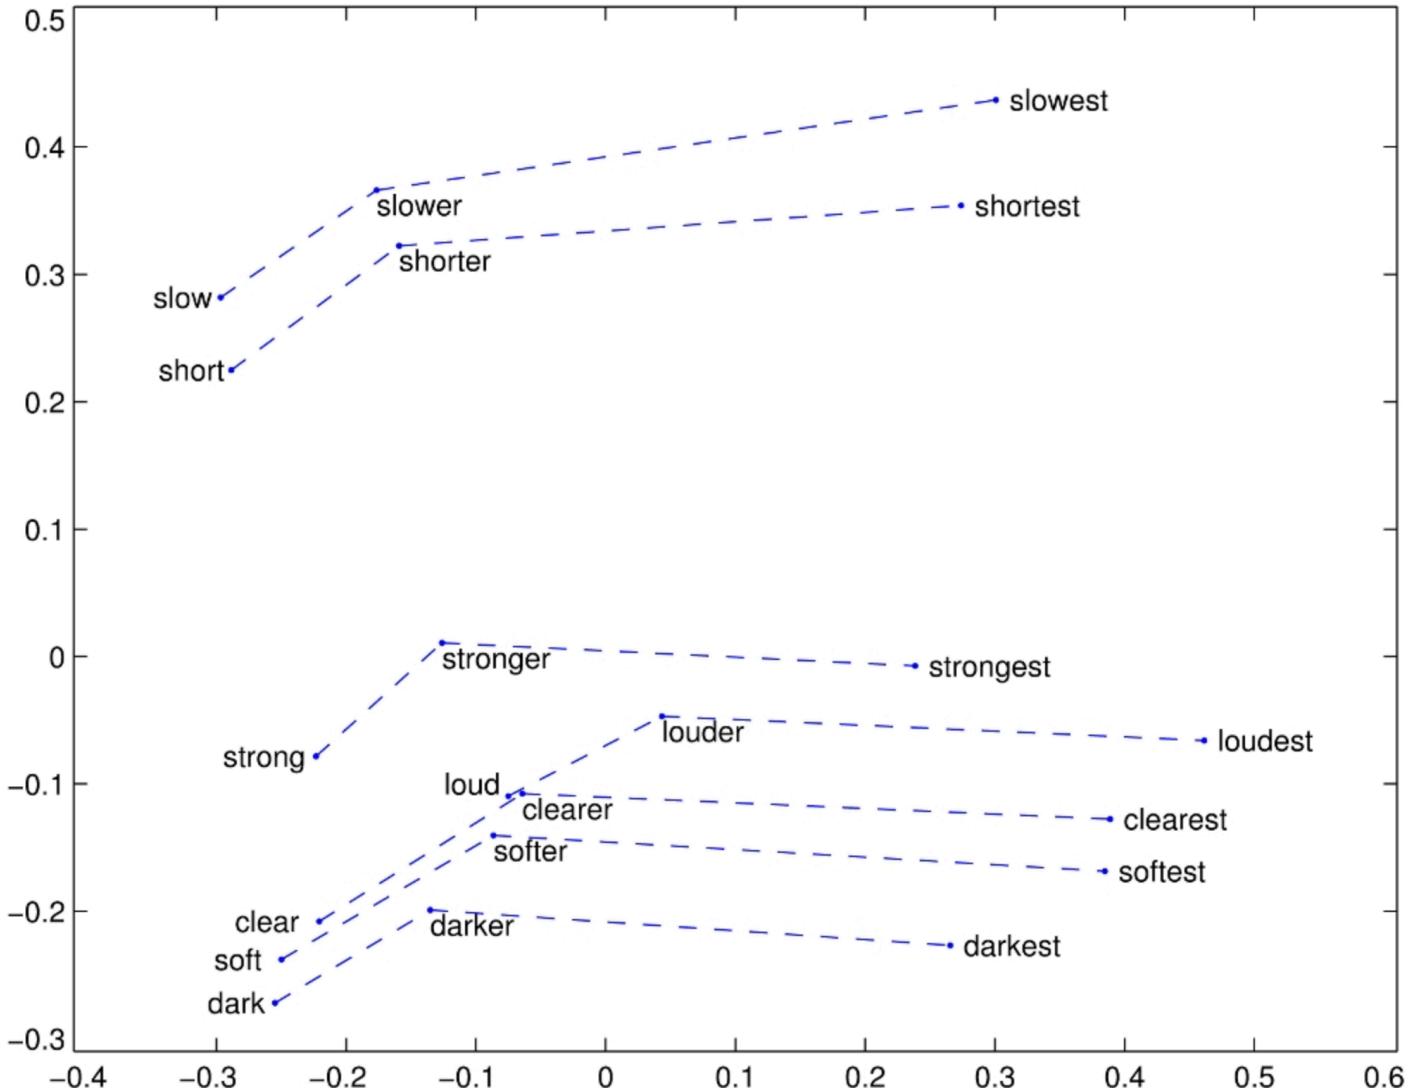
- Formula: 
$$\text{similarity} = \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}}$$

- For text, takes on a value between zero (totally unrelated) to one (meaning exactly terms appear in same proportion).
- Interpreted as angle between term vectors.
- Normalized for document length.

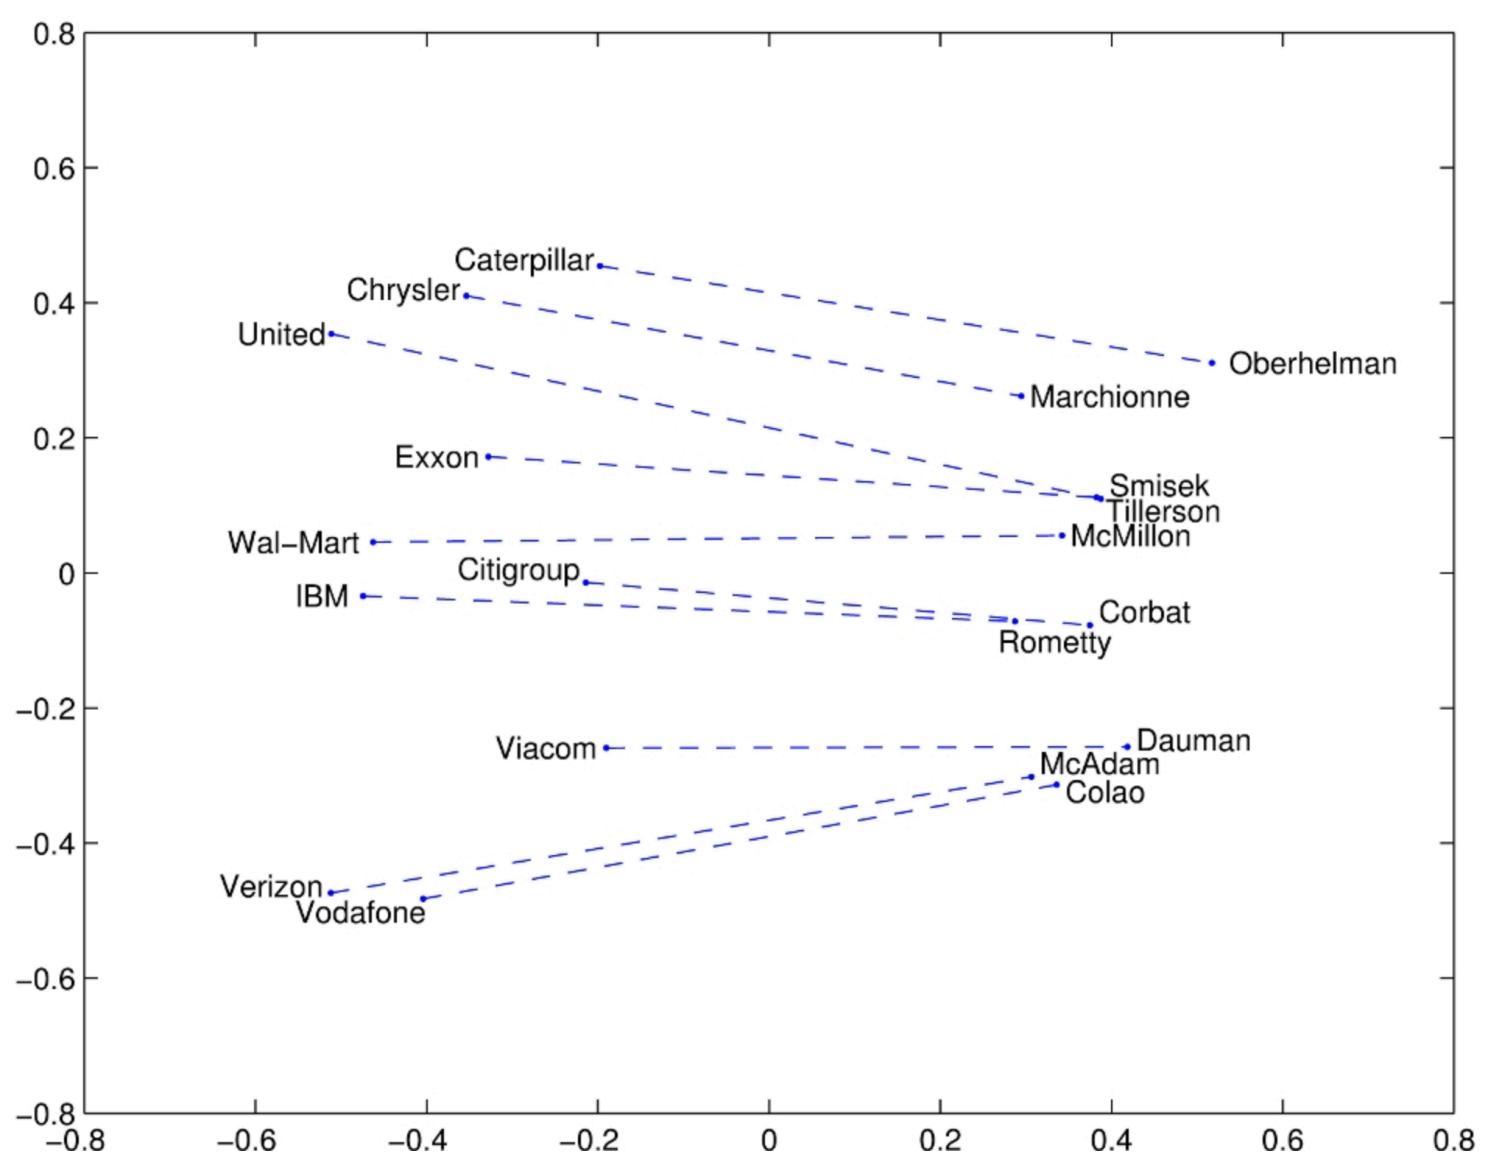
# Gender Relationships



# Superlative Relationships



# Company-CEO Relationships



# Solving Analogies

- **US** is to **Canada** as **hamburger** is to ??
- Relationships are encoded by word vector differences:  
$$f(\text{"canada"}) - f(\text{"us"}) = f(\text{"??"}) - f(\text{"hamburger"})$$
- We can add the relationship encoding vector to a the vector of **hamburger**:  
$$f(\text{"hamburger"}) + (f(\text{"canada"}) - f(\text{"us"})) = f(\text{"??"})$$

```
[('poutine', 0.4857853055000305),  
 ('cheeseburger', 0.4791874885559082),  
 ('hotdog', 0.47264325618743896),  
 ('australie', 0.43578580021858215),  
 ('croissant', 0.435305118560791),  
 ('hamburgers', 0.4350370168685913),  
 ('burger', 0.42210400104522705),  
 ('australien', 0.420423686504364),  
 ('abendblatt', 0.4082910120487213),  
 ('schweiz', 0.4046550989151001)]
```

# Solving Analogies

king – man + woman ~ queen



# Social Science Applications

- This is the tricky part...
- Train embeddings on different corpora to estimate differences in word association.
  - Gender and racial bias.
  - Over time.
  - Changes in language use.
- Document similarity, Author similarity/scaling.
- Preprocessing augmentation to combine two different sources of text.

# My Take

- Vector space representations of words will become increasingly important over the next few decades with the continued rise of deep learning methods.
- Right now, only a limited set of social science applications can really benefit from using them.
- In many cases, using pre-trained embeddings is the way to go. Need huge amounts of data and processing power to get excellent general purpose results.
- You need a lot more training to know what you are doing when fitting word embedding models. Good opportunity for collaboration.